

Problem Statement: To Predict how likely individuals are to receive their H1N1 and Seasonal Flu Vaccines.

Aim:

- To create and assess machine learning models that can precisely determine if individuals are likely to obtain the H1N1 and/or seasonal flu vaccines, utilizing demographic, behavioral, and health-related characteristics.

Dataset Overview:

- training_set_features.csv

Rows: 26,707

Columns: 36

Contains demographic, behavioral, and health-related survey responses.

- test_set_features.csv

Rows: 26,708

Columns: 36

Same structure as training features, without the vaccination labels.

- training_set_labels.csv

Rows: 26,707

Columns: 3

Contains vaccination labels for H1N1 and seasonal flu.

- The dataset consists of survey responses gathered to comprehend the factors affecting the adoption of H1N1 and seasonal flu vaccinations. It contains three primary files: features for training, labels for training, and features for testing. The training dataset includes 26,707 entries and 36 attributes, encompassing a diverse array of demographic, behavioral, health-related, and attitudinal factors.

Objective:

- Understand the dataset by performing exploratory data analysis (EDA) to identify patterns, distributions, and missing values.
- Handle missing data using appropriate imputation techniques, such as KNN imputation.
- Preprocess the data through encoding of categorical variables and normalization where required.

- Address class imbalance, particularly in the H1N1 vaccine target variable, using techniques like SMOTE.
- Build and compare multiple classification models, including:
 - Logistic Regression
 - Decision Tree
 - Random Forest
 - Gradient Boosting
 - Support Vector Machine (SVM)
 - Naive Bayes
- Evaluate model performance using appropriate metrics such as accuracy, precision, recall, F1-score, and AUC-ROC.
- Select the best-performing model for each vaccine type and interpret the results.
- Submit predictions to the competition platform (e.g., DrivenData) and analyze leaderboard performance.
- Draw conclusions and suggest recommendations based on findings, including public health insights or targeted outreach strategies.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: import io
%cd "C:\Users\coolr\Dropbox\PC\Desktop\Project"

C:\Users\coolr\Dropbox\PC\Desktop\Project
```

```
In [3]: vaccinetrain = pd.read_csv("training_set_features.csv")
vaccinetest = pd.read_csv("test_set_features.csv")
vaccinelabels = pd.read_csv("training_set_labels.csv")
```

```
In [4]: print(vaccinetrain.shape)
print(vaccinetest.shape)
print(vaccinelabels.shape)

(26707, 36)
(26708, 36)
(26707, 3)
```

```
In [5]: vaccinetrain.isnull().sum().sort_values(ascending=False)
```

```
Out[5]: employment_occupation      13470
employment_industry      13330
health_insurance      12274
income_poverty      4423
doctor_recc_h1n1      2160
doctor_recc_seasonal      2160
rent_or_own      2042
employment_status      1463
marital_status      1408
education      1407
chronic_med_condition      971
child_under_6_months      820
health_worker      804
opinion_seas_sick_from_vacc      537
opinion_seas_risk      514
opinion_seas_vacc_effective      462
opinion_h1n1_sick_from_vacc      395
opinion_h1n1_vacc_effective      391
opinion_h1n1_risk      388
household_children      249
household_adults      249
behavioral_avoidance      208
behavioral_touch_face      128
h1n1_knowledge      116
h1n1_concern      92
behavioral_large_gatherings      87
behavioral_outside_home      82
behavioral_antiviral_meds      71
behavioral_wash_hands      42
behavioral_face_mask      19
hhs_geo_region      0
census_msa      0
sex      0
race      0
age_group      0
respondent_id      0
dtype: int64
```

```
In [6]: vaccinetest.isnull().sum().sort_values(ascending=False)
```

```
Out[6]: employment_occupation      13426
employment_industry      13275
health_insurance      12228
income_poverty      4497
doctor_recc_h1n1      2160
doctor_recc_seasonal      2160
rent_or_own      2036
employment_status      1471
marital_status      1442
education      1407
chronic_med_condition      932
child_under_6_months      813
health_worker      789
opinion_seas_sick_from_vacc      521
opinion_seas_risk      499
opinion_seas_vacc_effective      452
opinion_h1n1_vacc_effective      398
opinion_h1n1_risk      380
opinion_h1n1_sick_from_vacc      375
household_children      225
household_adults      225
behavioral_avoidance      213
behavioral_touch_face      128
h1n1_knowledge      122
h1n1_concern      85
behavioral_outside_home      82
behavioral_antiviral_meds      79
behavioral_large_gatherings      72
behavioral_wash_hands      40
behavioral_face_mask      19
hhs_geo_region      0
census_msa      0
sex      0
race      0
age_group      0
respondent_id      0
dtype: int64
```

Columns:

```
In [7]: vaccinetrain.columns

# To know the column names in the dataset
```

```
Out[7]: Index(['respondent_id', 'h1n1_concern', 'h1n1_knowledge',
              'behavioral_antiviral_meds', 'behavioral_avoidance',
              'behavioral_face_mask', 'behavioral_wash_hands',
              'behavioral_large_gatherings', 'behavioral_outside_home',
              'behavioral_touch_face', 'doctor_recc_h1n1', 'doctor_recc_seasonal',
              'chronic_med_condition', 'child_under_6_months', 'health_worker',
              'health_insurance', 'opinion_h1n1_vacc_effective', 'opinion_h1n1_risk',
              'opinion_h1n1_sick_from_vacc', 'opinion_seas_vacc_effective',
              'opinion_seas_risk', 'opinion_seas_sick_from_vacc', 'age_group',
              'education', 'race', 'sex', 'income_poverty', 'marital_status',
              'rent_or_own', 'employment_status', 'hhs_geo_region', 'census_msa',
              'household_adults', 'household_children', 'employment_industry',
              'employment_occupation'],
              dtype='object')
```

```
In [8]: vaccinetest.columns

# To know the column names in the dataset
```

```
Out[8]: Index(['respondent_id', 'h1n1_concern', 'h1n1_knowledge',
              'behavioral_antiviral_meds', 'behavioral_avoidance',
              'behavioral_face_mask', 'behavioral_wash_hands',
              'behavioral_large_gatherings', 'behavioral_outside_home',
              'behavioral_touch_face', 'doctor_recc_h1n1', 'doctor_recc_seasonal',
              'chronic_med_condition', 'child_under_6_months', 'health_worker',
              'health_insurance', 'opinion_h1n1_vacc_effective', 'opinion_h1n1_risk',
              'opinion_h1n1_sick_from_vacc', 'opinion_seas_vacc_effective',
              'opinion_seas_risk', 'opinion_seas_sick_from_vacc', 'age_group',
              'education', 'race', 'sex', 'income_poverty', 'marital_status',
              'rent_or_own', 'employment_status', 'hhs_geo_region', 'census_msa',
              'household_adults', 'household_children', 'employment_industry',
              'employment_occupation'],
             dtype='object')
```

```
In [9]: vaccinelabels.columns

# To know the column names in the dataset
```

```
Out[9]: Index(['respondent_id', 'h1n1_vaccine', 'seasonal_vaccine'], dtype='object')
```

```
In [10]: # Since there are missing values in both Train and Test Set Im combining both the D
combinedf = pd.concat([vaccinetrain,vaccinetest],axis=0) # Row Concatenation
```

```
In [11]: # info() method shows some of the characteristics of the data such as Column Name,
#Dtype of the data, and Memory Usage.
combinedf.info()
# Many Variables have Null Values
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 53415 entries, 0 to 26707
Data columns (total 36 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   respondent_id                        53415 non-null  int64
1   h1n1_concern                        53238 non-null  float64
2   h1n1_knowledge                      53177 non-null  float64
3   behavioral_antiviral_meds           53265 non-null  float64
4   behavioral_avoidance                52994 non-null  float64
5   behavioral_face_mask                53377 non-null  float64
6   behavioral_wash_hands               53333 non-null  float64
7   behavioral_large_gatherings         53256 non-null  float64
8   behavioral_outside_home             53251 non-null  float64
9   behavioral_touch_face               53159 non-null  float64
10  doctor_recc_h1n1                   49095 non-null  float64
11  doctor_recc_seasonal               49095 non-null  float64
12  chronic_med_condition              51512 non-null  float64
13  child_under_6_months              51782 non-null  float64
14  health_worker                      51822 non-null  float64
15  health_insurance                   28913 non-null  float64
16  opinion_h1n1_vacc_effective         52626 non-null  float64
17  opinion_h1n1_risk                   52647 non-null  float64
18  opinion_h1n1_sick_from_vacc         52645 non-null  float64
19  opinion_seas_vacc_effective         52501 non-null  float64
20  opinion_seas_risk                   52402 non-null  float64
21  opinion_seas_sick_from_vacc         52357 non-null  float64
22  age_group                          53415 non-null  object
23  education                          50601 non-null  object
24  race                               53415 non-null  object
25  sex                                53415 non-null  object
26  income_poverty                     44495 non-null  object
27  marital_status                     50565 non-null  object
28  rent_or_own                        49337 non-null  object
29  employment_status                  50481 non-null  object
30  hhs_geo_region                     53415 non-null  object
31  census_msa                         53415 non-null  object
32  household_adults                   52941 non-null  float64
33  household_children                 52941 non-null  float64
34  employment_industry                26810 non-null  object
35  employment_occupation              26519 non-null  object
dtypes: float64(23), int64(1), object(12)
memory usage: 15.1+ MB
```

Missing Value Imputation: Strategy (KNN Imputer)

- KNN is a non parametric algorithm that uses distance like Euclidean to calculate distance between missing value and closest 5 neighbors(rows) and smallest distance value is imputed.
- Before applying KNN imputer labelencoding of data must be done

```
In [12]: from sklearn.preprocessing import LabelEncoder
```

```
C:\Users\coolr\anaconda3\lib\site-packages\scipy\__init__.py:155: UserWarning: A NumPy version >=1.18.5 and <1.26.0 is required for this version of SciPy (detected version 1.26.4
  warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

```
In [13]: stringcols=combinedf.drop(['respondent_id'],axis=1)
```

```
In [14]: stringcols=stringcols.astype(str).apply(lambda series:pd.Series(
LabelEncoder().fit_transform(series[series.notnull()]),index = series[series.notnu]

In [15]: from sklearn.impute import KNNImputer

In [16]: imputer = KNNImputer()

In [17]: stringcolsimp = imputer.fit_transform(stringcols)

In [18]: stringcolsimp = pd.DataFrame(stringcolsimp,columns=stringcols.columns)

In [19]: stringcolsimp.head()
```

```
Out[19]:
```

| | h1n1_concern | h1n1_knowledge | behavioral_antiviral_meds | behavioral_avoidance | behavioral_fac |
|---|--------------|----------------|---------------------------|----------------------|----------------|
| 0 | 1.0 | 0.0 | 0.0 | 0.0 | |
| 1 | 3.0 | 2.0 | 0.0 | 1.0 | |
| 2 | 1.0 | 1.0 | 0.0 | 1.0 | |
| 3 | 1.0 | 1.0 | 0.0 | 1.0 | |
| 4 | 2.0 | 1.0 | 0.0 | 1.0 | |

5 rows × 35 columns

Splitting back the data to Train data and Test data

```
In [20]: vaccinetraindf = stringcolsimp.loc[0:26706]

In [21]: vaccinetestdf = stringcolsimp.loc[26708:53415]

In [22]: vaccinetraindf.shape

Out[22]: (26707, 35)

In [23]: vaccinetestdf.shape

Out[23]: (26707, 35)
```

Hypothesis Testing:

a.) Seasonal Vaccine

```
In [24]: from scipy.stats import chi2_contingency

In [25]: # Null Hypothesis : There is Significance between seasonalvaccine and sex
pd.crosstab(vaccinelabels.seasonal_vaccine,vaccinetraindf.sex)
```

Out[25]:

| | sex | 0.0 | 1.0 |
|------------------|------|------|-----|
| seasonal_vaccine | | | |
| 0 | 7970 | 6302 | |
| 1 | 7888 | 4547 | |

In [26]: `chi2_contingency(pd.crosstab(vaccinelabels.seasonal_vaccine,vaccinetraindf.sex))`
Since p value < 0.05 reject null

Out[26]: (158.40462897107128,
2.5248076319253385e-36,
1,
array([[8474.38409406, 5797.61590594],
[7383.61590594, 5051.38409406]]))

In [27]: *# Null Hypothesis : There is Significance between seasonalvaccine and income_poverty*
`pd.crosstab(vaccinelabels.seasonal_vaccine,vaccinetraindf.income_poverty)`

Out[27]:

| | income_poverty | 0.0 | 1.0 | 2.0 | 3.0 |
|------------------|----------------|------|------|------|-----|
| seasonal_vaccine | | | | | |
| 0 | 6686 | 3427 | 1719 | 2440 | |
| 1 | 6091 | 3383 | 978 | 1983 | |

In [28]: `chi2_contingency(pd.crosstab(vaccinelabels.seasonal_vaccine,vaccinetraindf.income_p`
Since p value < 0.05 reject null

Out[28]: (153.17012658009583,
5.455880245363045e-33,
3,
array([[6827.92316621, 3639.20769836, 1441.25450256, 2363.61463287],
[5949.07683379, 3170.79230164, 1255.74549744, 2059.38536713]]))

In [29]: *# Null Hypothesis : There is Significance between seasonalvaccine and education*
`pd.crosstab(vaccinelabels.seasonal_vaccine,vaccinetraindf.education)`

Out[29]:

| | education | 0.0 | 1.0 | 2.0 | 3.0 | 4.0 |
|------------------|-----------|------|------|------|-----|-----|
| seasonal_vaccine | | | | | | |
| 0 | 3199 | 1415 | 4933 | 3860 | 865 | |
| 1 | 2598 | 948 | 5164 | 3183 | 542 | |

In [30]: `chi2_contingency(pd.crosstab(vaccinelabels.seasonal_vaccine,vaccinetraindf.educatio`
Since p value < 0.05 reject null

Out[30]: (173.57815735793127,
1.7841137139002192e-36,
4,
array([[3097.86887333, 1262.76766391, 5395.7533231 , 3763.72097203,
751.88916763],
[2699.13112667, 1100.23233609, 4701.2466769 , 3279.27902797,
655.11083237]]))

In [31]: *# Null Hypothesis : There is Significance between seasonalvaccine and doctor_recc_s*
`pd.crosstab(vaccinelabels.seasonal_vaccine,vaccinetraindf.doctor_recc_seasonal)`


```
Out[31]: doctor_recc_seasonal    0.0    1.0    2.0
```

| seasonal_vaccine | | | | |
|------------------|------|-------|------|------|
| | 0 | 10756 | 2117 | 1399 |
| 1 | 5697 | 5977 | 761 | |

```
In [32]: chi2_contingency(pd.crosstab(vaccinelabels.seasonal_vaccine,vaccinetraindf.doctor_r  
# Since p value < 0.05 reject null
```

```
Out[32]: (3474.902794649153,  
0.0,  
2,  
array([[8792.3471749 , 4325.36668289, 1154.28614221],  
[7660.6528251 , 3768.63331711, 1005.71385779]]))
```

```
In [33]: # Null Hypothesis : There is Significance between seasonalvaccine and behavioral_fa  
pd.crosstab(vaccinelabels.seasonal_vaccine,vaccinetraindf.behavioral_face_mask)
```

```
Out[33]: behavioral_face_mask    0.0    1.0    2.0
```

| seasonal_vaccine | | | | |
|------------------|-------|-------|-----|---|
| | 0 | 13449 | 815 | 8 |
| 1 | 11398 | 1026 | 11 | |

```
In [34]: chi2_contingency(pd.crosstab(vaccinelabels.seasonal_vaccine,vaccinetraindf.behavior  
# Since p value < 0.05 reject null
```

```
Out[34]: (67.92302367431634,  
1.7811595018273916e-15,  
2,  
array([[1.32780314e+04, 9.83815180e+02, 1.01534429e+01],  
[1.15689686e+04, 8.57184820e+02, 8.84655708e+00]]))
```

b.) H1N1_ Vaccine

```
In [35]: # Null Hypothesis : There is no significance between h1n1vaccine and sex  
pd.crosstab(vaccinelabels.h1n1_vaccine,vaccinetraindf.sex)
```

```
Out[35]: sex    0.0    1.0
```

| h1n1_vaccine | | |
|--------------|------|------------|
| | 0 | 12378 8655 |
| 1 | 3480 | 2194 |

```
In [36]: chi2_contingency(pd.crosstab(vaccinelabels.h1n1_vaccine,vaccinetraindf.sex))  
# Since p value < 0.05 reject null
```

```
Out[36]: (11.309981152455997,  
0.0007709155489949327,  
1,  
array([[12488.90979893, 8544.09020107],  
[ 3369.09020107, 2304.90979893]]))
```

```
In [37]: # Null Hypothesis : There is no significance between seasonalvaccine and income_pov  
pd.crosstab(vaccinelabels.h1n1_vaccine,vaccinetraindf.income_poverty)
```

```
Out[37]: income_poverty    0.0    1.0    2.0    3.0
```

| h1n1_vaccine | |
|--------------|----------------------|
| 0 | 10178 5087 2181 3587 |
| 1 | 2599 1723 516 836 |

```
In [38]: chi2_contingency(pd.crosstab(vaccinelabels.h1n1_vaccine,vaccinetraindf.income_poverty))
# Since p value < 0.05 reject null
```

```
Out[38]: (94.91226392182662,
1.9282818704757377e-20,
3,
array([[10062.4795372 , 5363.19054929, 2124.01246864, 3483.31744486],
       [ 2714.5204628 , 1446.80945071, 572.98753136, 939.68255514]]))
```

```
In [39]: # Null Hypothesis : There is no significance between h1n1_vaccine and education
pd.crosstab(vaccinelabels.h1n1_vaccine,vaccinetraindf.education)
```

```
Out[39]: education    0.0    1.0    2.0    3.0    4.0
```

| h1n1_vaccine | | | | | | |
|--------------|---|------|------|------|------|------|
| | 0 | 4726 | 1968 | 7614 | 5579 | 1146 |
| | 1 | 1071 | 395 | 2483 | 1464 | 261 |

```
In [40]: chi2_contingency(pd.crosstab(vaccinelabels.h1n1_vaccine,vaccinetraindf.education))
# Since p value < 0.05 reject null
```

```
Out[40]: (130.12342530584624,
3.66443514946588e-27,
4,
array([[4565.40611076, 1860.97199236, 7951.85535627, 5546.68884562,
       1108.07769499],
       [1231.59388924, 502.02800764, 2145.14464373, 1496.31115438,
       298.92230501]]))
```

```
In [41]: # Null Hypothesis : There is no significance between h1n1_vaccine and doctor_recc_seasonal
pd.crosstab(vaccinelabels.h1n1_vaccine,vaccinetraindf.doctor_recc_seasonal)
```

```
Out[41]: doctor_recc_seasonal    0.0    1.0    2.0
```

| h1n1_vaccine | |
|--------------|-----------------|
| 0 | 13783 5275 1975 |
| 1 | 2670 2819 185 |

```
In [42]: chi2_contingency(pd.crosstab(vaccinelabels.h1n1_vaccine,vaccinetraindf.doctor_recc_seasonal))
# Since p value < 0.05 reject null
```

```
Out[42]: (1347.6234370031846,
2.329639991754564e-293,
2,
array([[12957.49986895, 6374.40004493, 1701.10008612],
       [ 3495.50013105, 1719.59995507, 458.89991388]]))
```

```
In [43]: # Null Hypothesis : There is no significance between h1n1_vaccine and behavioral_face_mask
pd.crosstab(vaccinelabels.h1n1_vaccine,vaccinetraindf.behavioral_face_mask)
```

```
Out[43]: behavioral_face_mask    0.0    1.0    2.0
```

| h1n1_vaccine | | | | |
|--------------|---|-------|------|----|
| | 0 | 19765 | 1255 | 13 |
| | 1 | 5082 | 586 | 6 |


```
In [44]: chi2_contingency(pd.crosstab(vaccinelabels.h1n1_vaccine,vaccinetraindf.behavioral_f  
# Since p value < 0.05 reject null
```

```
Out[44]: (133.81661811498145,  
8.751662937660443e-30,  
2,  
array([[1.95681638e+04, 1.44987280e+03, 1.49633804e+01],  
[5.27883619e+03, 3.91127195e+02, 4.03661961e+00]]))
```

```
In [45]: y1 = vaccinelabels.h1n1_vaccine  
y = vaccinelabels.seasonal_vaccine  
X = vaccinetraindf
```


```
In [46]: vaccinelabels.seasonal_vaccine.plot(kind='hist')  
# Balanced Data
```

```
Out[46]: <AxesSubplot: ylabel='Frequency'>
```



```
In [47]: vaccinelabels.h1n1_vaccine.plot(kind='hist')  
# Imbalance Data
```

```
Out[47]: <AxesSubplot: ylabel='Frequency'>
```



Methods

Model Building and Ideology :

- As there are two binary variables to target, Let's create a classification model for each variable separately.
- The variable seasonal vaccine is balanced, whereas the h1n1 vaccine variable is not.
- Let's use SMOTE to balance the h1n1 vaccine variable.

```
In [48]: from sklearn.metrics import classification_report
from sklearn.model_selection import cross_val_score
from sklearn.metrics import plot_roc_curve
```

Logistic Regression for Seasonal Vaccine

```
In [49]: from sklearn.linear_model import LogisticRegression
```

```
In [50]: logit = LogisticRegression(max_iter=300)
```

```
In [51]: logitmodel = logit.fit(X,y)
```

```
In [52]: logitmodel.score(X,y)
```

```
Out[52]: 0.7610738757629086
```

```
In [53]: cross_val_score(logitmodel,X,y)
```

```
Out[53]: array([0.75739423, 0.75870461, 0.76427635, 0.76652312, 0.75360419])
```

```
In [54]: np.mean([0.75739423, 0.759079 , 0.76427635, 0.76671035, 0.75360419])
```

```
Out[54]: 0.760212824
```

```
In [55]: logitpredict = logitmodel.predict(X)
```

```
In [56]: logitpredict_test = logitmodel.predict_proba(vaccinetestdf)
```

```
In [57]: pd.DataFrame(logitpredict_test).to_csv("seaslog.csv")
```

```
In [58]: print(classification_report(y,logitpredict))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.77 | 0.79 | 0.78 | 14272 |
| 1 | 0.75 | 0.73 | 0.74 | 12435 |
| accuracy | | | 0.76 | 26707 |
| macro avg | 0.76 | 0.76 | 0.76 | 26707 |
| weighted avg | 0.76 | 0.76 | 0.76 | 26707 |

```
In [59]: pd.crosstab(y,logitpredict)
```

```
Out[59]:
```


| | col_0 | 0 | 1 |
|------------------|-------|------|---|
| seasonal_vaccine | | | |
| 0 | 11277 | 2995 | |
| 1 | 3386 | 9049 | |

```
In [60]: plot_roc_curve(logitmodel,X,y)
```

C:\Users\coolr\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: Future Warning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

warnings.warn(msg, category=FutureWarning)

```
Out[60]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x19883229820>
```



Decision Tree Classification for Seasonal Vaccine

```
In [61]: from sklearn.tree import DecisionTreeClassifier
```

```
In [62]: tree = DecisionTreeClassifier(max_depth=16)
```

```
In [63]: treemodel = tree.fit(X,y)
```

```
In [64]: treemodel.score(X,y)
```

```
Out[64]: 0.9402029430486389
```

```
In [65]: cross_val_score(treemodel,X,y)
```

```
Out[65]: array([0.69842756, 0.70254586, 0.70717094, 0.71241341, 0.69930725])
```

```
In [66]: np.mean([0.6963684 , 0.70647697, 0.70642202, 0.71241341, 0.69668601])
```

```
Out[66]: 0.7036733620000001
```

```
In [67]: treepredict = treemodel.predict(X)
```

```
In [68]: treepredict_test = treemodel.predict(vaccinetestdf)
```

```
In [69]: pd.DataFrame(treepredict_test).to_csv("seastree.csv")
```

```
In [70]: print(classification_report(y,treepredict))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.94 | 0.95 | 0.94 | 14272 |
| 1 | 0.94 | 0.93 | 0.94 | 12435 |
| accuracy | | | 0.94 | 26707 |
| macro avg | 0.94 | 0.94 | 0.94 | 26707 |
| weighted avg | 0.94 | 0.94 | 0.94 | 26707 |

```
In [71]: pd.crosstab(y,treepredict)
```

```
Out[71]:
```


| | col_0 | 0 | 1 |
|------------------|-------|-------|---|
| seasonal_vaccine | | | |
| 0 | 13508 | 764 | |
| 1 | 833 | 11602 | |

```
In [72]: plot_roc_curve(treemodel,X,y)
```

C:\Users\coolr\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: Future Warning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

warnings.warn(msg, category=FutureWarning)

```
Out[72]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x198833c7eb0>
```



Random Forest Classification for Seasonal Vaccine

```
In [73]: from sklearn.ensemble import RandomForestClassifier
```

```
In [74]: RF = RandomForestClassifier(n_estimators=3)
```

```
In [75]: RFmodel = RF.fit(X,y)
```

```
In [76]: RFmodel.score(X,y)
```

```
Out[76]: 0.9460441082862171
```

```
In [77]: cross_val_score(RFmodel,X,y)
```

```
Out[77]: array([0.71602396, 0.71246724, 0.72477064, 0.71934095, 0.70941771])
```

```
In [78]: np.mean([0.70497941, 0.70778735, 0.71316233, 0.72477064, 0.71035387])
```

```
Out[78]: 0.7122107200000001
```

```
In [79]: RFpredict = RFmodel.predict(X)
```

```
In [80]: RFpredict_test = RFmodel.predict_proba(vaccinetestdf)
```

```
In [81]: pd.DataFrame(RFpredict_test).to_csv('RFseas.csv')
```

```
In [82]: print(classification_report(y,RFpredict))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.95 | 0.95 | 0.95 | 14272 |
| 1 | 0.94 | 0.94 | 0.94 | 12435 |
| accuracy | | | 0.95 | 26707 |
| macro avg | 0.95 | 0.95 | 0.95 | 26707 |
| weighted avg | 0.95 | 0.95 | 0.95 | 26707 |

```
In [83]: pd.crosstab(y,RFpredict)
```

```
Out[83]:
```


| | col_0 | 0 | 1 |
|------------------|-------|-------|---|
| seasonal_vaccine | | | |
| 0 | 13589 | 683 | |
| 1 | 758 | 11677 | |

```
In [84]: plot_roc_curve(RFmodel,X,y)
```

C:\Users\coolr\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

warnings.warn(msg, category=FutureWarning)

```
Out[84]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x198fda4dd60>
```

Gradient Boosting Classification for Seasonal Vaccine

```
In [85]: from sklearn.ensemble import GradientBoostingClassifier
```

```
In [86]: gbm = GradientBoostingClassifier(n_estimators=4000)
```

```
In [87]: gbmmodel = gbm.fit(X,y)
```

```
In [88]: gbmmodel.score(X,y)
```

```
Out[88]: 0.8566667914778897
```

```
In [89]: cross_val_score(gbmmodel,X,y)
```

```
Out[89]: array([0.77031075, 0.77162112, 0.77588467, 0.78168882, 0.77813144])
```

```
In [90]: np.mean([0.77405466, 0.77461625, 0.77925482, 0.78674406, 0.78468452])
```

```
Out[90]: 0.7798708619999999
```

```
In [91]: gbmpredict = gbmmodel.predict(X)
```

```
In [92]: gbmpredict_test = gbmmodel.predict_proba(vaccinetestdf)
```

```
In [93]: pd.DataFrame(gbmpredict_test).to_csv('gbmseas.csv')
```

```
In [94]: print(classification_report(y,gbmpredict))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.86 | 0.87 | 0.87 | 14272 |
| 1 | 0.85 | 0.84 | 0.84 | 12435 |
| accuracy | | | 0.86 | 26707 |
| macro avg | 0.86 | 0.86 | 0.86 | 26707 |
| weighted avg | 0.86 | 0.86 | 0.86 | 26707 |

```
In [95]: pd.crosstab(y,gbmpredict)
```

```
Out[95]:
```


| | col_0 | 0 | 1 |
|------------------|-------|-------|---|
| seasonal_vaccine | | | |
| 0 | 12460 | 1812 | |
| 1 | 2016 | 10419 | |

```
In [96]: plot_roc_curve(gbmmodel,X,y)
```

C:\Users\coolr\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

warnings.warn(msg, category=FutureWarning)

```
Out[96]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x198eb60c160>
```



Support Vector Machine Classification for Seasonal Vaccine

```
In [97]: from sklearn.svm import SVC
```

```
In [98]: svm = SVC()
```

```
In [99]: svmmodel = svm.fit(X,y)
```

```
In [100]: svmmodel.score(X,y)
```

```
Out[100]: 0.7783352679072902
```

```
In [101]: cross_val_score(svmmodel,X,y)
```

```
Out[101]: array([0.76843879, 0.76506926, 0.77682082, 0.77625913, 0.77195282])
```

```
In [102]: np.mean([0.76843879, 0.76506926, 0.77682082, 0.77625913, 0.77195282])
```

```
Out[102]: 0.7717081640000001
```

```
In [103]: svmpredict_test = svmmodel.predict(X)
```

```
In [104]: pd.DataFrame(svmpredict_test).to_csv('svmseas.csv')
```

```
In [105]: print(classification_report(y,svmpredict_test))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.78 | 0.81 | 0.80 | 14272 |
| 1 | 0.77 | 0.74 | 0.76 | 12435 |
| accuracy | | | 0.78 | 26707 |
| macro avg | 0.78 | 0.78 | 0.78 | 26707 |
| weighted avg | 0.78 | 0.78 | 0.78 | 26707 |

```
In [106]: pd.crosstab(y,svmpredict_test)
```

```
Out[106]:
```


| | col_0 | 0 | 1 |
|------------------|-------|------|---|
| seasonal_vaccine | | | |
| 0 | 11580 | 2692 | |
| 1 | 3228 | 9207 | |

```
In [107]: plot_roc_curve(svmmodel,X,y)
```

C:\Users\coolr\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: Future Warning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

warnings.warn(msg, category=FutureWarning)

```
Out[107]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x19885d5ca30>
```



Navie Bayes Classification for Seasonal Vaccine

```
In [108...] from sklearn.naive_bayes import BernoulliNB
```

```
In [109...] nb = BernoulliNB()
```

```
In [110...] nbmodel = nb.fit(X,y)
```

```
In [111...] nbmodel.score(X,y)
```

```
Out[111]: 0.6936383719624069
```

```
In [112...] cross_val_score(nbmodel,X,y)
```

```
Out[112]: array([0.68045676, 0.70423063, 0.69106909, 0.69855832, 0.68938401])
```

```
In [113...] np.mean([0.68045676, 0.70423063, 0.69106909, 0.69855832, 0.68938401])
```

```
Out[113]: 0.692739762
```

```
In [114...] nbpredict = nbmodel.predict(X)
```

```
In [115...] nbpredict_test = nbmodel.predict_proba(vaccinetestdf)
```

```
In [116...] pd.DataFrame(nbpredict_test).to_csv('nbseas.csv')
```

```
In [117...] print(classification_report(y,nbpredict))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.72 | 0.69 | 0.71 | 14272 |
| 1 | 0.66 | 0.70 | 0.68 | 12435 |
| accuracy | | | 0.69 | 26707 |
| macro avg | 0.69 | 0.69 | 0.69 | 26707 |
| weighted avg | 0.69 | 0.69 | 0.69 | 26707 |

```
In [118... pd.crosstab(y,nbpredict)
```

```
Out[118]:
```


| | col_0 | 0 | 1 |
|------------------|-------|------|---|
| seasonal_vaccine | | | |
| 0 | 9878 | 4394 | |
| 1 | 3788 | 8647 | |

```
In [119... plot_roc_curve(nbmodel,X,y)
```

C:\Users\coolr\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: Future Warning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

warnings.warn(msg, category=FutureWarning)

```
Out[119]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x19885e78a90>
```



Balancing the Imbalance data using Smote

```
In [120... from imblearn.over_sampling import SMOTE
```

```
In [121...] smote = SMOTE(random_state=60)

In [122...] X_smote,y_smote = smote.fit_resample(X,y1)

In [123...] print(pd.DataFrame(y1).value_counts())
print(pd.DataFrame(y_smote).value_counts())
```

```
h1n1_vaccine
0          21033
1          5674
dtype: int64
h1n1_vaccine
0          21033
1          21033
dtype: int64
```

Logistic Regression for H1N1 Vaccine

```
In [124...] print(X_smote.shape)
print(y_smote.shape)
```

```
(42066, 35)
(42066,)
```

```
In [125...] logit2 = LogisticRegression(max_iter=400)
```

```
In [126...] logitmodel_smote = logit2.fit(X_smote,y_smote)
```

```
In [127...] logitmodel_smote.score(X_smote,y_smote)
```

```
Out[127]: 0.762943945228926
```

```
In [128...] cross_val_score(logitmodel_smote,X_smote,y_smote)
```

```
Out[128]: array([0.75112907, 0.75680494, 0.76726495, 0.77011768, 0.76667063])
```

```
In [129...] np.mean([0.75124792, 0.75692381, 0.76714608, 0.77011768, 0.76678949])
```

```
Out[129]: 0.7624449959999999
```

```
In [130...] logitmodel_smote.predict(X_smote)
```

```
In [131...] logitmodel_smote.predict_proba(vaccinetestdf)
```

```
In [132...] pd.DataFrame(logitmodel_smote.predict_test).to_csv('h1n1log.csv')
```

```
In [133...] print(classification_report(y_smote,logitmodel_smote.predict))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.76 | 0.76 | 0.76 | 21033 |
| 1 | 0.76 | 0.77 | 0.76 | 21033 |
| accuracy | | | 0.76 | 42066 |
| macro avg | 0.76 | 0.76 | 0.76 | 42066 |
| weighted avg | 0.76 | 0.76 | 0.76 | 42066 |

```
In [134... pd.crosstab(y_smote,logitsmote_predict)
```

```
Out[134]:
```


| | col_0 | 0 | 1 |
|--------------|-------|-------|---|
| h1n1_vaccine | | | |
| 0 | 15990 | 5043 | |
| 1 | 4929 | 16104 | |

```
In [135... plot_roc_curve(logitmodel_smote,X_smote,y_smote)
```

C:\Users\coolr\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: Future Warning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

warnings.warn(msg, category=FutureWarning)

```
Out[135]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x19885c69e80>
```



Decision Tree Classification for H1N1 Vaccine :

```
In [136... treesmote = DecisionTreeClassifier(max_depth=11)
```

```
In [137... treemodel_smote = treesmote.fit(X_smote,y_smote)
```

```
In [138... treemodel_smote.score(X_smote,y_smote)
```

```
Out[138]: 0.9075262682451386
```

```
In [139... cross_val_score(treemodel_smote,X_smote,y_smote)
```

```
Out[139]: array([0.69562634, 0.85379769, 0.93022703, 0.93640794, 0.93153453])
```

```
In [140... np.mean([0.69550749, 0.85391656, 0.93010817, 0.9365268 , 0.93212885])
```

```
Out[140]: 0.8696375740000001
```

```
In [141... treepredict_smote = treemodel_smote.predict(X_smote)
```

```
In [142... treepredict_smote_test =treemodel_smote.predict_proba(vaccinetestdf)
```

```
In [143... pd.DataFrame(treepredict_smote_test).to_csv('h1n1tree.csv')
```

```
In [144... print(classification_report(y_smote,treepredict_smote))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.88 | 0.95 | 0.91 | 21033 |
| 1 | 0.94 | 0.87 | 0.90 | 21033 |
| accuracy | | | 0.91 | 42066 |
| macro avg | 0.91 | 0.91 | 0.91 | 42066 |
| weighted avg | 0.91 | 0.91 | 0.91 | 42066 |

```
In [145... pd.crosstab(y_smote,treepredict_smote)
```

```
Out[145]:
```


| | col_0 | 0 | 1 |
|---------------------|-------|-------|---|
| h1n1_vaccine | | | |
| 0 | 19880 | 1153 | |
| 1 | 2737 | 18296 | |

```
In [146... plot_roc_curve(treemodel_smote,X_smote,y_smote)
```

C:\Users\coolr\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: Future Warning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

```
warnings.warn(msg, category=FutureWarning)
```

```
Out[146]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x19885f37a60>
```

Random Forest Classification for H1N1 Vaccine :

```
In [147...] RFsmote = RandomForestClassifier(n_estimators=7,max_depth=12)
```

```
In [148...] RFmodel_smote = RFsmote.fit(X_smote,y_smote)
```

```
In [149...] RFmodel_smote.score(X_smote,y_smote)
```

```
Out[149]: 0.9265440022821281
```

```
In [150...] cross_val_score(RFmodel_smote,X_smote,y_smote)
```

```
Out[150]: array([0.69681483, 0.88006656, 0.95245453, 0.95209794, 0.94449067])
```

```
In [151...] np.mean([0.6931305 , 0.87388565, 0.94591703, 0.95304885, 0.94413408])
```

```
Out[151]: 0.8820232220000002
```

```
In [152...] RFsmote_predict = RFmodel_smote.predict(X_smote)
```

```
In [153...] RFsmote_predict_test = RFmodel_smote.predict(vaccinetestdf)
```

```
In [154...] pd.DataFrame(RFsmote_predict_test).to_csv('h1h1RF.csv')
```

```
In [155...] print(classification_report(y_smote,RFsmote_predict))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.91 | 0.95 | 0.93 | 21033 |
| 1 | 0.95 | 0.90 | 0.92 | 21033 |
| accuracy | | | 0.93 | 42066 |
| macro avg | 0.93 | 0.93 | 0.93 | 42066 |
| weighted avg | 0.93 | 0.93 | 0.93 | 42066 |

```
In [156... pd.crosstab(y_smote, RFsmote_predict)
```

```
Out[156]:
```


| | col_0 | 0 | 1 |
|---------------------|-------|-------|---|
| h1n1_vaccine | | | |
| 0 | 20021 | 1012 | |
| 1 | 2078 | 18955 | |

```
In [157... plot_roc_curve(RFmodel_smote, X_smote, y_smote)
```

C:\Users\coolr\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: Future Warning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

warnings.warn(msg, category=FutureWarning)

```
Out[157]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x19885f82e50>
```



Gradient Boosting Classification for H1N1 Vaccine :

```
In [158... gbm_smote = GradientBoostingClassifier(n_estimators=700)
```

```
In [159... gbmmmodel_smote = gbm_smote.fit(X_smote,y_smote)
```

```
In [160... gbmmmodel_smote.score(X_smote,y_smote)
```

```
Out[160]: 0.9202443778823753
```

```
In [161... cross_val_score(gbmmmodel_smote,X_smote,y_smote)
```

```
Out[161]: array([0.64879962, 0.88458338, 0.96576726, 0.96529181, 0.96208249])
```

```
In [162... np.mean([0.64606608, 0.8851777 , 0.96612386, 0.9656484 , 0.9630334 ])
```

```
Out[162]: 0.885209888
```

```
In [163... gbmpredict_smote = gbmmmodel_smote.predict(X_smote)
```

```
In [164... gbmpredict_smote_test = gbmmmodel_smote.predict_proba(vaccinetestdf)
```

```
In [165... pd.DataFrame(gbmpredict_smote_test).to_csv('h1n1gbm.csv')
```

```
In [166... print(classification_report(y_smote,gbmpredict_smote))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.89 | 0.96 | 0.92 | 21033 |
| 1 | 0.95 | 0.88 | 0.92 | 21033 |
| accuracy | | | 0.92 | 42066 |
| macro avg | 0.92 | 0.92 | 0.92 | 42066 |
| weighted avg | 0.92 | 0.92 | 0.92 | 42066 |

```
In [167... pd.crosstab(y_smote,gbmpredict_smote)
```

```
Out[167]:
```


| | col_0 | 0 | 1 |
|---------------------|-------|-------|---|
| h1n1_vaccine | | | |
| 0 | 20099 | 934 | |
| 1 | 2421 | 18612 | |

```
In [168... plot_roc_curve(gbmmmodel_smote,X_smote,y_smote)
```

C:\Users\coolr\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: Future Warning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

```
warnings.warn(msg, category=FutureWarning)
```

```
Out[168]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x1988621aeb0>
```



Support Vector Classification for H1N1 Vaccine :

```
In [169...] svmmodel_smote = svm.fit(X_smote,y_smote)
```

```
In [170...] svmmodel_smote.score(X_smote,y_smote)
```

```
Out[170]: 0.8064470118385394
```

```
In [171...] cross_val_score(svmmodel_smote,X_smote,y_smote)
```

```
Out[171]: array([0.75861659, 0.79044336, 0.81267087, 0.81564246, 0.80387496])
```

```
In [172...] np.mean([0.75861659, 0.79044336, 0.81267087, 0.81564246, 0.80387496])
```

```
Out[172]: 0.796249648
```

```
In [173...] svmsmote_predict = svmmodel_smote.predict(X_smote)
```

```
In [174...] svmsmote_predict_test = svmmodel_smote.predict(vaccinetestdf)
```

```
In [175...] pd.DataFrame(svmsmote_predict_test).to_csv('h1n1svm.csv')
```

```
In [176...] print(classification_report(y_smote,gbmpredict_smote))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.89 | 0.96 | 0.92 | 21033 |
| 1 | 0.95 | 0.88 | 0.92 | 21033 |
| accuracy | | | 0.92 | 42066 |
| macro avg | 0.92 | 0.92 | 0.92 | 42066 |
| weighted avg | 0.92 | 0.92 | 0.92 | 42066 |

```
In [177... pd.crosstab(y_smote,gbmpredict_smote)
```

```
Out[177]:
```


| | col_0 | 0 | 1 |
|---------------------|-------|-------|---|
| h1n1_vaccine | | | |
| 0 | 20099 | 934 | |
| 1 | 2421 | 18612 | |

```
In [178... plot_roc_curve(svmmodel_smote,X_smote,y_smote)
```

C:\Users\coolr\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: Future Warning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

warnings.warn(msg, category=FutureWarning)

```
Out[178]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x19886279310>
```



Naive Bayes Classification for H1N1 Vaccine :

```
In [179... nbmodel_smote = nb.fit(X_smote,y_smote)
```

```
In [180...] nbmodel_smote.score(X_smote,y_smote)
```

```
Out[180]: 0.7696001521418723
```

```
In [181...] cross_val_score(nbmodel_smote,X_smote,y_smote)
```

```
Out[181]: array([0.64606608, 0.75121835, 0.81397837, 0.81195768, 0.80744087])
```

```
In [182...] np.mean([0.64606608, 0.75121835, 0.81397837, 0.81195768, 0.80744087])
```

```
Out[182]: 0.76613227
```

```
In [183...] nbmote_predict = nbmodel_smote.predict(X_smote)
```

```
In [184...] nbmote_predict_test = nbmodel_smote.predict_proba(vaccinetestdf)
```

```
In [185...] pd.DataFrame(nbmote_predict_test).to_csv('h1n1nb.csv')
```

```
In [186...] print(classification_report(y_smote,nbmote_predict))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.78 | 0.75 | 0.76 | 21033 |
| 1 | 0.76 | 0.79 | 0.77 | 21033 |
| accuracy | | | 0.77 | 42066 |
| macro avg | 0.77 | 0.77 | 0.77 | 42066 |
| weighted avg | 0.77 | 0.77 | 0.77 | 42066 |

```
In [187...] pd.crosstab(y_smote,nbmote_predict)
```

```
Out[187]:
```

| | col_0 | 0 | 1 |
|---------------------|-------|-------|---|
| h1n1_vaccine | | | |
| 0 | 15716 | 5317 | |
| 1 | 4375 | 16658 | |


```
In [188...] from sklearn.metrics import RocCurveDisplay
```

```
In [189...] plot_roc_curve(nbmodel_smote,X_smote,y_smote)
```

C:\Users\coolr\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: Future Warning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

warnings.warn(msg, category=FutureWarning)

```
Out[189]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x198858f15b0>
```



Results

Performance Metric Report:

| Seasonal Vaccine | | | | | | | |
|----------------------------------|-----------------|----------|---|-----------|--------|----------|------|
| Model | Cross Val Score | Accuracy | | Precesion | Recall | F1 Score | AUC |
| Logistic Regression | 76.02 | 76.11 | 0 | 0.77 | 0.79 | 0.78 | 0.83 |
| | | | 1 | 0.75 | 0.73 | 0.74 | |
| Decision Tree Classification | 70.36 | 94.05 | 0 | 0.94 | 0.95 | 0.94 | 0.99 |
| | | | 1 | 0.94 | 0.93 | 0.94 | |
| Random Forest Classification | 71.22 | 94.6 | 0 | 0.95 | 0.95 | 0.95 | 0.98 |
| | | | 1 | 0.95 | 0.94 | 0.94 | |
| Support Vector Classification | 77.17 | 77.83 | 0 | 0.78 | 0.81 | 0.80 | 0.85 |
| | | | 1 | 0.77 | 0.74 | 0.76 | |
| Naive Bayes | 69.27 | 69.36 | 0 | 0.72 | 0.69 | 0.71 | 0.76 |
| | | | 1 | 0.66 | 0.7 | 0.68 | |
| Gradient Boosting Classification | 77.98 | 85.66 | 0 | 0.86 | 0.87 | 0.87 | 0.93 |
| | | | 1 | 0.85 | 0.84 | 0.84 | |

| H1N1 Vaccine | | | | | | | |
|----------------------------------|-----------------|----------|---|-----------|--------|----------|------|
| Model | Cross Val Score | Accuracy | | Precesion | Recall | F1 Score | AUC |
| Logistic Regression | 76.24 | 76.29 | 0 | 0.76 | 0.76 | 0.76 | 0.84 |
| | | | 1 | 0.76 | 0.77 | 0.76 | |
| Decision Tree Classification | 86.96 | 90.75 | 0 | 0.88 | 0.95 | 0.91 | 0.97 |
| | | | 1 | 0.94 | 0.87 | 0.90 | |
| Random Forest Classification | 88.20 | 92.42 | 0 | 0.90 | 0.95 | 0.93 | 0.98 |
| | | | 1 | 0.95 | 0.9 | 0.92 | |
| Support Vector Classification | 79.62 | 80.64 | 0 | 0.89 | 0.96 | 0.92 | 0.88 |
| | | | 1 | 0.95 | 0.88 | 0.92 | |
| Naive Bayes | 76.61 | 76.96 | 0 | 0.78 | 0.75 | 0.76 | 0.85 |
| | | | 1 | 0.76 | 0.79 | 0.77 | |
| Gradient Boosting Classification | 88.52 | 92.02 | 0 | 0.89 | 0.96 | 0.92 | 0.97 |
| | | | 1 | 0.95 | 0.88 | 0.92 | |

Key Observation:

Seasonal Vaccine prediction

- Gradient Boosting achieved the best balance of performance, with accuracy of 85.66% and AUC of 0.93.
- Random Forest and Decision Tree models also performed well, with very high recall and precision (AUCs of 0.98 and 0.99 respectively), though Decision Tree may be overfitting given its accuracy of 94.05% on a lower cross-validation score (70.36).
- Naive Bayes and Logistic Regression showed weaker generalization with lower accuracy and F1 scores.

H1N1 Vaccine prediction

- Gradient Boosting again outperformed other models with accuracy of 92.02% and AUC of 0.97.
- Random Forest and Decision Tree closely followed, both showing high predictive power (AUC > 0.97).
- Support Vector Classification had strong recall (0.96 for class 0), which is valuable in public health applications where identifying individuals unlikely to take the vaccine is important.
- Naive Bayes and Logistic Regression underperformed compared to ensemble methods.

Conclusion:

- For both vaccine types, Gradient Boosting Classification proved to be the most reliable and robust model. It achieved high scores across all evaluation metrics while balancing between precision and recall, making it highly suitable for public health prediction tasks. This model could help healthcare providers target outreach more effectively to individuals less likely to get vaccinated.

References:

Aim ref:

- <https://www.drivendata.org/competitions/66/flu-shot-learning/>

Objective ref:

- <https://www.cdc.gov/mmwr/volumes/68/rr/rr6803a1.htm>

Data set I took from this website:

- <https://www.drivendata.org/competitions/66/flu-shot-learning/data/>